
Unsupervised Cost-Sensitive Online Prediction

Anonymous Author(s)

Affiliation

Address

email

Abstract

Abstract goes here

1 Introduction

The quality of sensors used in a decision system influences the accuracy of the measurements or the predictions. Typically, a less accurate sensor is cheap and produces the predictions faster. A more accurate sensor is costly and takes more time to output predictions. In practice, the budget/time constraints do not allow costly/slow sensors to be used all the time and one has to use the sensor that is the most ‘cost effective’. One natural way is to use a sensor for which sum of prediction error rate and cost is the lowest. However, values of the errors rate may not be known a priori and the best cost effective sensor cannot be determined. Further, the true labels required to estimate the error rate may also be not available or prohibitively expensive to know. In this paper we study the problem of learning the most cost effective sensor when the true labels are not known, and the goal is to efficiently learn the most cost effective sensor.

In summary

1. *there is a test-time acquisition problem. This is not us. Because we do not have annotated training data.*

2. *This is not an online version of test-time acquisition problem. Because we do not get ground truth labels for the predictions we make.*

3. *We are in a situation where we have a network of sensors and their predictions if and when we probe those sensors. Our situation is such that we observe the predictions of the sensors in the directed neighborhood of the probed sensor.*

4. *Thus we do not directly observe the label and the situation resembles a online/bandit setting but without observing either a noisy or noiseless version of the reward.*

5. *This situation arises in Homeland security as well as in medical diagnosis.*

This problem arises in many applications including homeland security, communication networks and medical diagnosis. For example, in the homeland security problem, where bags need to be screened for potential threats, either a cheap Infra-red (IR) imager or a more expensive and time consuming active millimeter wave (AMMW) scanner can be used. In medical diagnosis, practitioners can use non-invasive blood test, CT scan to determine a medical condition or go for a more invasive surgical procedures. In wireless communications, network designer can use error correcting codes of different block lengths to overcome channel noise. A code with higher block length (more redundancy) improve the tolerance against noise but reduces transmission rate.

Several papers including Trapeznikov & Saligrama (2013), Trapeznikov et al. (2014) Xu et al. (2013) have considered the problem of learning the best cost effective predictor/classifier using supervised learning methods. The general approach in these methods is to learn a decision function by minimizing an empirical risk objective over a training set. The objective functions in these methods are inherently non-convex and the authors resort to convex relaxations and experimental validations without any theoretical guarantees. However, in many applications gathering training samples may be infeasible,

and moreover the labels may not be available at all. We consider an online version of this problem where the samples arrive sequentially and a learner has to decide which sensors to apply for prediction. For each sample, the learner only observes sensor predictions and true label is not revealed.

In this work we focus on sequential predication of binary labels. Similar to Trapeznikov et al. (2014), we consider that the order in which sensors are applied is fixed. Typically, the cheapest sensor, or the one with highest error rate, is used first, followed by next cheap sensor with smaller error rate and so on. The sensors thus constitute stages of a cascade, where prediction error rates decrease along the depth, while the costs increase. For each new sample, the learner applies the sensors sequentially and can stop at any stage in the cascade. The goal is to stop at a stage where expected loss is the smallest. Loss at depth k is defined as total cost incurred for acquiring sensor predictions plus a penalty which is 1 if the prediction of k^{th} sensor is correct, otherwise it is zero. If the learner stops at a depth k , he obtains the predictions of all the first k sensors as feedback, but which of them are correct is unknown. We refer to this setup as the Sensor Acquisitions Problem (SAP).

The feedback in SAP do not reveal information about the losses, hence the learner cannot identify the best stage for any sample. We thus focus on scenarios where feedback satisfies some stochastic ordering. Specifically, we assume that if a sensor in the cascade predicts a label correctly, any subsequent sensor also predict it correctly. We refer to this assumption as *dominance condition*. When it holds, the learner can partially infer losses of the stages, which, as discussed later, is sufficient to learn the best stage for a given sample. We further demonstrate that under any weaker condition the learner cannot identify the best stage. Dominance conditions holds in many scenarios includes the examples discussed at the beginning. In the wireless communication example, if an error correcting code (ex. Reed-Solomon, LDPC Mackay recovers information in a channel with certain noise level, then with more redundancy blocks in the error correction code we can certainly recover the information on the channel (though at a lower transmission rate).

Our first main contribution is to show that if the dominance condition holds the SAP problem can be reduced to a stochastic multi-armed bandit with side observations, where bandit arms are identified with the stage of cascade, the payoff of an arm is given by loss from the corresponding stage, and side observation structure is defined by the feedback graph induced by the cascade. In particular, we show that the SAPregret of any meta-strategy is equal to its bandit-regret when the procedure is used to play in the corresponding bandit problem. As a consequence, we conclude that existing efficient bandit algorithms, as well as their bounds on bandit-regret, can be directly applied to achieve new results for SAP. Although the underlying reduction is straightforward, it gives ready policies with performance guarantees for SAP and their fundamental limitations .

Related Work: Trapeznikov & Saligrama (2013)Seldin et al. (2014)

Structure of paper

2 Sensor Acquisition Problem

The learner has access to $K \geq 2$ sensors that are ordered in terms of their prediction efficiency. Specifically, we consider that the sensors form a cascade (order in which the sensors are selected is predetermined) and in each round the learner can sequentially select a subset of sensors in the cascade and stop at any depth.

Let $\{Z_t, Y_t\}_{t \geq 0}$ denote a sequence generated according to an unknown distribution. $Z_t \in \mathcal{C} \subset \mathcal{R}^d$, where \mathcal{C} is a compact set, denotes a feature vector/context at time t and $Y_t \in \{0, 1\}$ its binary label. We denote output/prediction of the i^{th} sensor as \hat{Y}_t^i when its input is Z_t . The set of actions available to the learner is $\mathcal{A} = \{1, \dots, K\}$, where the action $k \in \mathcal{A}$ indicates acquiring predictions from sensors $1, \dots, k$ and classifying using the prediction \hat{Y}_t^k .

The prediction error rate of the i^{th} sensor is denoted as $\gamma_i := \Pr\{Y_t \neq \hat{Y}_t^i\}$. In this section we assume that the error rate does not depend on the context and postpone the treatment with contextual information to Section 6. Further, the sensors are arranged such that the prediction error rate improves with depth in the cascade, i.e., $\gamma_{k-1} \geq \gamma_k$ for all $k > 2$. However, the learner incurs an extra cost of $c_k \geq 0$ to acquire output of




Figure 1: Cascade of sensors

90 sensor k after acquiring output of sensor $k - 1$. The sensor
91 cascade is depicted in the adjacent figure.

92 Let $H_t(k)$ denote the feedback observed in round t from action
93 k . Since we observe predictions of all the first k sensors by
94 playing action k , we get $H_t(k) = \{\hat{Y}_t^1, \dots, \hat{Y}_t^k\}$. The loss
95 incurred in each round is defined in terms of the prediction error and the total cost involved. When
96 the learner selects action k , loss is the prediction error of sensor k plus sum of the costs incurred
97 along the path (c_1, \dots, c_k) . Let $L_t : \mathcal{A} \rightarrow \mathcal{R}_+$ denote the loss function in round t . Then,

$$L_t(k) = \mathbf{1}_{\{\hat{Y}_t^k \neq Y_t\}} + \sum_{j=1}^k c_j. \quad (1)$$

98 We refer to the above setup as Sensor Acquisition Problem (SAP) and denote it as $\psi =$
99 $(K, \mathcal{A}, (\gamma_i, c_{i-1})_{i \in [K]})^1$. A policy $\pi^\psi = (\pi_1^\psi, \pi_2^\psi, \dots)$ on ψ , where $\pi_t^\psi : \mathcal{H}_{t-1} \rightarrow \mathcal{A}$, gives ac-
100 tion selected in each round using history \mathcal{H}_{t-1} that consists of all actions and corresponding feedback
101 observed before t . Let Π^ψ denote set of policies on ψ . For any $\pi \in \Pi^\psi$, we compare its performance
102 with respect to the optimal policy (single best action in hindsight) and define its expected regret as
103 follows

$$R_T^\psi(\pi) = \mathbb{E} \left[\sum_{t=1}^T L_t(a_t) \right] - \min_{k \in \mathcal{A}} \mathbb{E} \left[\sum_{t=1}^T L_t(k) \right], \quad (2)$$

104 where a_t denotes the policy selected by π_t in round t . The goal of the learner is to learn a policy that
105 minimizes the expected total loss, or, equivalently, to minimize the expected regret, i.e.,

$$\pi^* = \arg \min_{\pi \in \Pi^\psi} R_T^\psi(\pi). \quad (3)$$

106 **Optimal action in hindsight:** For any t , we have

$$\mathbb{E}[L_t(k)] = \Pr\{Y_t \neq \hat{Y}_t^k\} + \sum_{j=1}^k c_j = \gamma_k + \sum_{j=1}^k c_j. \quad (4)$$

107 Let $k^* = \arg \min_{k \in \mathcal{A}} \gamma_k + \sum_{i < k} c_i$. Then the optimal policy is to play action k^* in each round. If an
108 action i is played in any round then it adds $\Delta_k := \gamma_k + \sum_{i < k} c_i - (\gamma_{k^*} + \sum_{i < k^*} c_i)$ to the expected
109 regret. Let I_t denote the action selected in round t and $N_k^\psi(s)$ denote the number of times action k is
110 selected till time s , i.e., $N_k^\psi(s) = \sum_{t=1}^s \mathbf{1}_{\{I_t=k\}}$. Then the expected regret can be expressed as

$$R_T^\psi(\pi) = \sum_{k \in \mathcal{A}} \mathbb{E}[N_k^\psi(T)] \Delta_k. \quad (5)$$

111

112 3 When is SAP Learnable?

113 In the SA-Problem feedback $H_t(\cdot)$ does not reveal any information about the true label Y_t in any
114 round t . Hence the loss values are not known, and we are in a hopeless situation where linear regret
115 is unavoidable. In this section we explore conditions that lead to policies that are Hannan consistent
116 Hanna (1957), i.e, a policy $\pi \in \Pi^\psi$ such that $R_T^\psi(\pi)/T \rightarrow 0$.

117 Let us consider $K = 2$ sensors and start with a simple condition that if sensor 1 predicts the label 1
118 correctly, then sensor 2 also predicts it correctly², i.e.,

$$Y_t = 1 \text{ and } \hat{Y}_t^1 = 1 \implies \hat{Y}_t^2 = 1. \quad (6)$$

119 To fix ideas, we enumerate all the 8 possible tuples $(Y, \hat{Y}^1, \hat{Y}^2)$ as shown in Table 3, and write
120 probability of the i th tuple $i = 1, 2, \dots, 8$ as p_{i-1} . From Table 3, we have $\gamma_1 = p_2 + p_3 + p_4 + p_5$
121 and $\gamma_2 = p_1 + p_3 + p_4 + p_6$, thus

$$\gamma_1 - \gamma_2 = p_2 + p_5 - p_1 - p_6. \quad (7)$$

¹Note that $k \in \mathcal{A}$ implies that action k selects all sensors $1, 2, \dots, k$, not just sensor k . We set $c_0 = 0$

²Suppose we interpret label 1 as 'threat', the condition implies that if sensor 1 detects threat correctly, the better sensor 2 also detects it.

Y	\hat{Y}_t^1	\hat{Y}_t^2	$\Pr(Y, \hat{Y}^1, \hat{Y}^2)$
0	0	0	p_0
0	0	1	p_1
0	1	0	p_2
0	1	1	p_3
1	0	0	p_4
1	0	1	p_5
1	1	0	p_6
1	1	1	p_7

$$\Pr(\hat{Y}^1, \hat{Y}^2) = \begin{cases} p_1 + p_5 & \text{if } (\hat{Y}^1, \hat{Y}^2) = (0, 1) \\ p_2 + p_6 & \text{if } (\hat{Y}^1, \hat{Y}^2) = (1, 0) \\ p_0 + p_4 & \text{if } (\hat{Y}^1, \hat{Y}^2) = (0, 0) \\ p_3 + p_7 & \text{if } (\hat{Y}^1, \hat{Y}^2) = (1, 1) \end{cases} \quad (8)$$

From (4), action 1 is optimal if $\gamma_1 - \gamma_2 \leq c$, otherwise action 2 is optimal. If a policy learns the difference $\gamma_1 - \gamma_2$, it can play the optimal arm and it is Hannan consistent. Note that only sensor output (\hat{Y}^1, \hat{Y}^2) are observed and not the true label Y . Hence only values of marginal probabilities $\Pr(\hat{Y}^1, \hat{Y}^2)$ as given in (8) can be used to learn the difference $\gamma_1 - \gamma_2$. The following example demonstrate that just knowing the values of marginals is not enough to decide which action is optimal.

Set $c = 0.35$ and consider the following two case: 1) $p_2 = 1/2, p_1 = 1/4 - 1/40, p_5 = 1/4 + 1/40$ and 2) $p_2 = 1/2, p_1 = 1/4 - 3/40, p_5 = 1/4 + 3/40$. From condition (6) we have $p_6 = 0$. Also, set $p_0 = p_4 = p_3 = p_7 = 0$ in both the cases. We get $\gamma_1 - \gamma_2 = 0.3$ in the first case, hence action 1 is optimal. Where as $\gamma_1 - \gamma_2 = 0.4$ in the second case, hence actions 2 is optimal. However, for both the cases the marginals $\Pr(\hat{Y}^1, \hat{Y}^2)$ are the same for all pairs (\hat{Y}^1, \hat{Y}^2) . Since we only observe the pairs (\hat{Y}^1, \hat{Y}^2) , one cannot hope to distinguish the cases and linear regret is unavoidable.

Next, assume that if sensor 0 predicts the label 0 correctly, then sensor 2 also predicts it correctly, i.e.,

$$Y_t = 0 \text{ and } \hat{Y}_t^1 = 0 \implies \hat{Y}_t^2 = 0. \quad (9)$$

We can argue similar to the previous example that under this conditions one cannot expect better than linear regret. Now assume that both (6) and (9) hold. Then, $p_2 = p_6 = 0$ and we get $\gamma_1 - \gamma_2 = p_5 - p_1$. Since $p_5 = \Pr(0, 1)$ and $p_1 = \Pr(1, 0)$, we can learn their values by observing $(0, 1)$ and $(1, 0)$ patterns and thus hope for a Hannan consistent policy. In the following we assume that (6) and (9) hold and refer to it as dominance condition. For the case of $K > 2$ sensors, it is given as follows:

Assumption 1 (Dominance Condition) *If sensor i predicts correctly, all the sensors in the subsequent stages of the cascade also predict correctly, i.e.,*

$$\hat{Y}_t^i = Y_t \rightarrow \hat{Y}_t^j \quad \forall j > i \geq 1 \quad (10)$$

In the following we establish that under the dominance condition efficient algorithms for a SAP problem can be derived from algorithms on a suitable stochastic multi-armed bandit problem. We first recall the stochastic multi-armed bandit setting and the relevant results.

4 Background on Stochastic Multi-armed Bandits

A stochastic multi-armed bandit (MAB), denoted as $\phi := (K, (\nu_k)_{1 \leq k \leq K})$, is a sequential learning problem where number of arms K is known and each arm $i \in [K]$ gives rewards drawn according to an unknown distribution ν_k . Let $X_{i,n}$ denote the random reward from arm i in its n th play. For each arm $i \in [K]$, $\{X_{i,t} : t > 0\}$ are independently and identically (i.i.d) distributed and for all $t > 0$, $\{X_{i,t}, i \in [K]\}$ are independent. We note that in the standard MAB setting the learner observes only reward from the selected arm in each round and no information from the other arms is revealed. A policy is any allocation strategy that maps the past history into an arm in each round, and let Π^ϕ denote a set of policies on ϕ . If the learner knows $\{\nu_k\}_{k \in [K]}$, then the optimal policy is to play the arm with highest mean. For any policy $\pi \in \Pi^\phi$, its performance is measured with respect to the optimal policy and is defined in terms of expected cumulative regret (or simply regret) as follows:

157 Let π selects arm i_t in round t . After T rounds, its regret is

$$R_T^\phi(\pi) = T\mu_{i^*} - \sum_{t=1}^T \mu_{i_t}, \quad (11)$$

158 where $\mu_i = \mathbb{E}[X_{i,n}]$ denotes mean of distribution ν_i for all $i \in [K]$ and $i^* = \arg \max_{i \in [K]} \mu_i$. Let
 159 $N_i^\phi(t) = \sum_{s=1}^t \mathbf{1}\{i_s = i\}$ denote the number of pulls of arm i till time t . Then, the Regret of policy
 160 π can be expressed

$$R_T^\phi(\pi) = \sum_{i=1}^K (\mu_{i^*} - \mu_i) \mathbb{E}[N_i^\phi(T)].$$

161 The goal of the learner is to learn a policy that minimizes the regret.

162 MAB problems have been extensively studied in the literature. The seminal paper of Lai & Robbins
 163 Lai & Robbins (1985) showed that for any consistent policy (that plays sub-optimal arms only
 164 sup-polynomially many times in the time horizon) the regret grows logarithmically in time horizon.
 165 Specifically, for a class of parametric reward distributions, they showed that regret of any consistent
 166 policy satisfies

$$\liminf_{n \rightarrow \infty} \frac{R_T^\phi(\pi)}{\log T} \geq \sum_{i \neq i^*} \frac{\mu_{i^*} - \mu_i}{D(\mu_{i^*} || \mu_i)}, \quad (12)$$

167 where $D(p||q)$ is the KL-divergence of $p, q \in [0, 1]$. Further, the authors in Lai & Robbins (1985)
 168 provided an upper confidence bound (UCB) based policy that asymptotically achieves the lower
 169 bound for a class of parametric reward distributions.

170 Auer et. al. Auer et al. (2002) proposed an anytime policy named UCB1, that is based on the UCB
 171 strategy and showed that it is optimal on any MAB with bounded rewards. Specifically, they showed
 172 that regret of UCB1 for any $T > K$ is upper bound as

$$R_T^\phi(\text{UCB1}) \leq \sum_{i \neq i^*} \frac{8 \log n}{\mu_{i^*} - \mu_i} + (\pi^2/3 + 1)(\mu_{i^*} - \mu_i). \quad (13)$$

173 Thus demonstrating the optimality of UCB1. Since the work of Auer et. al. several works have
 174 proposed improvised UCB based policies like, KL-UCB Garivier & Cappé (2011), MOSS Audibert
 175 & Bubeck (2010).

176 4.1 MAB With Side Information

177 In many applications playing an arm reveals information about the other arms which can be exploited
 178 to improve learning performance. Let \mathcal{N}_i denote the set of arms such that playing arm i reveals
 179 rewards of all arms $j \in \mathcal{N}_i$. We refer to \mathcal{N}_i as neighborhood of i and the graph induced by the
 180 neighborhood sets as side-information graph. Given a set of neighborhood $\{\mathcal{N}_i, i \in [K]\}$, let
 181 $\phi_G := (K, (\nu_k)_{1 \leq k \leq K}, G)$ denote a MAB with side-information graph $G = (V, E)$, where $|V| = K$
 182 and $(i, j) \in E$ if $j \in \mathcal{N}_i$. The side-observation graph is known to the learner and remains fixed
 183 during the play.

184 Let Π^{ϕ_G} denote the set of all policies on ϕ_G that map the past history (including the side-observations)
 185 to an action in each round. For any policy $\pi \in \Pi^{\phi_G}$, we denote the regret over a period T as $R_T^{\phi_G}(\pi)$
 186 and is given by (11). Note that, in each round, only reward from the arm played contribute to the
 187 regret and not that from the side-observations. In Buccapatnam et al. (2014) the authors extended
 188 the lower bound in (12) to incorporate the effect of side-observations. Specifically, they establish
 189 that any policy $\pi \in \Pi^{\phi_G}$ where side observation graph is such that $i \in \mathcal{N}_i$ for all $i \in [K]$ satisfies
 190 Buccapatnam et al. (2014)

$$\liminf_{T \rightarrow \infty} R_T^{\phi_G}(\pi) / \log T \geq \eta(G) \quad (14)$$

191 where $\eta(G)$ is the optimal value of the following linear program

$$\begin{aligned}
LP1 : \min_{\{w_i\}} & \sum_{i \in [K]} (\mu_{i^*} - \mu_i) w_i \\
\text{subjected to} & \sum_{j \in \mathcal{N}_i} w_j \geq 1/D(\mu_i || \mu_{i^*}) \text{ for all } i \in [K] \\
& w_i \geq 0 \text{ for all } i \in [K]
\end{aligned} \tag{15}$$

192 **Definition 1 (Domination number Buccapatnam et al. (2014))** *Given a graph $G = (V, E)$, a sub-*
193 *set $W \subset V$ is a dominant set if for each $v \in V$ there exists $u \in W$ such that $(u, v) \in E$. The size of*
194 *the smallest dominant set is called weak domination number and is denoted as $\xi(G)$.*

195 The authors in Buccapatnam et al. (2014) gave an UCB based strategy, named UCB-LP, that exploits
196 the side-observations and explore arms at a rate in proportion to the size of their neighborhood.
197 UCB-LP plays arms in proportions to the values $\{z_i^*, i \in [K]\}$ computed from the following linear
198 programmer which is a relaxation of linear programme $LP1$.

$$\begin{aligned}
LP2 : \min_{\{z_i\}} & \sum_{i \in [K]} z_i \\
\text{subjected to} & \sum_{j \in \mathcal{N}_i} z_j \geq 1 \text{ for all } i \in [K] \\
& z_i \geq 0 \text{ for all } i \in [K]
\end{aligned} \tag{16}$$

199 The regret of UCB-LP is upper bounded by

$$\mathcal{O} \left(\sum_{i \in [K]} z_i^* \log T \right) + \mathcal{O}(K\delta), \tag{17}$$

200 where $\delta = \max_{i \in [K]} |\mathcal{K}_i|$ and $\{z_i^*\}$ are the optimal values of $LP2$. Since any dominating set is a
201 feasible solution of $LP2$, we get $\sum_{i \in [K]} z_i^* \leq \xi(G)$, and the regret of UCB-LP is $\mathcal{O}(\xi(G) \log T)$.

202 4.2 Special case: 1-armed bandit

203 In the MAB problem when all the arms have a fixed reward except for one, we get 1-armed bandit.
204 The learner knows the arms that give fixed reward the goal is to identify the quality of the arm that
205 gives stochastic reward as fast as possible. A straightforward modification of UCB1 achieves optimal
206 regret of $\Theta(\log T)$.

207 5 Regret Equivalence

208 In this section we establish that under the dominance condition SAP is ‘regret equivalent’ to an
209 instance of MAB with side-information and the corresponding algorithm for MAB can be suitably
210 imported to solve SAP efficiently.

211 **Definition 2 (Regret Equivalence)** *Consider a SAP problem $\psi := (K, \mathcal{A}, (\gamma_i, c_{i-1})_{i \in [K]})$ and a*
212 *bandit problem with $\phi_G := (N, (\nu_i)_{i \in [N]}, G)$ side-information graph G . We say that ψ is regret-*
213 *equivalent to ϕ_G if given a policy π for problem ψ , one can come up with a policy π' that uses π ,*
214 *such that the regret of π' on any instance of ϕ_G is the same as the regret of π on some corresponding*
215 *instance of ψ , and vice versa.*

216 In the following we first consider the SAP with 2 sensors and then the general case with more than 2
217 sensors. The 2 sensors case helps to draw comparison with the well studied apple tasting problem
218 and understand role of the dominance condition.

219 5.1 SAP with two sensors

220 In the SAP with only two actions, the feedback from action $i = 1$ reveals no information about
221 the loss incurred in that round. However feedback after action $i = 2$ reveals (partial) information

about the loss of both actions. Suppose feedback is such that predictions of the sensors disagree, i.e., $\hat{Y}_t^1 \neq \hat{Y}_t^2$ after action 2. The dominance condition then implies that the only possible events are $\hat{Y}_t^1 \neq Y_t$ and $\hat{Y}_t^2 = Y_t$. I.e., the true label is that predicted by sensor-2, hence loss incurred is just c (prediction loss is zero). Suppose predictions of the sensors agree, i.e., $\hat{Y}_t^1 = \hat{Y}_t^2$, then the dominance condition implies that either predictions of both are correct or both are incorrect. Though the true loss is not known in this case, the learner can infer some useful knowledge: in round t , if prediction of both the sensors agree, then the difference in losses of the actions is $L_t(2) - L_t(1) = c > 0$. And if predictions of the sensors disagree, then dominance assumption implies that $L_t(1) = 1$ and $L_t(2) = c$ or $L_t(2) - L_t(1) = c - 1 < 0$. Thus, each time learner plays action 2, he gets to know whether or not he was better off by selecting the other action. This setup sounds similar to the standard apple tasting problem Helmboet et al. (2000)], but differs in terms of the information structure when action 2 is played.

Apple tasting problem: In the apple tasting problem, a learner gets a sequence of apples and some of them can be rotten. In each round, the learner can either accept or reject an apple. If an apple is accepted, the learner tastes it and incurs a penalty if it is rotten. If apple is rejected, he still incurs the penalty if it is rotten, but do not get to observe its quality. The goal of the learner is to taste more good apples. The SAP setting is a more general version than the apple tasting problem—in any round, actions 1 reveals no loss values. Action 2 reveals only partial information about the losses, but not the exact losses as in the apple tasting problem. However, we next show that the partial information is enough to achieve optimal performance.

Theorem 1 Assume dominance condition (10) holds. Then SAP ψ with $K = 2$ is regret-equivalent to a stochastic 1-armed bandit.

The following corollary follow immediately from the regret equivalence.

Proposition 1 (SAP regret lower bound) Let π be any policy on SAT with 2 sensors such that it pulls the suboptimal arm only sub polynomial many times, i.e., $\mathbb{E}[N_i(T)] = o(T^a)$ for all $a > 0$ and $i \neq i^*$. Then,

$$\liminf_{T \rightarrow \infty} R_T^\psi(\pi) / \log T \geq \frac{|\gamma_1 - \gamma_2 - c|}{D(\hat{\gamma}, \gamma^*)} \text{ where } \gamma^* = \min\{\gamma_1, \gamma_2 + c\}, \hat{\gamma} = \max\{\gamma_1, \gamma_2 + c\} \quad (18)$$

and $D(\hat{\gamma}, \gamma^*)$ is the KL-divergence between $\hat{\gamma}$ and γ^* .

Proposition 2 (SAP regret upper bound) Let π' denote a policy on a 1-armed stochastic bandit where one arm has mean $\gamma_1 - \gamma_2$ and the other gives fixed reward c . Then, the regret of a policy $g(\pi)$ for the SAT problem obtained according the mapping (27) is upper bounded as

$$R_T^\psi(g(\pi)) \leq \frac{6 \log T}{|\gamma_1 - \gamma_2 - c|} + |\gamma_1 - \gamma_2 - c|(1 + \pi^2/3) \text{ when } \pi' = \text{UCB1}. \quad (19)$$

$$R_T^\psi(g(\pi)) \leq \frac{|\gamma_1 - \gamma_2 - c| \log T}{D(\hat{\gamma}, \gamma^*)} + \mathcal{O}(\sqrt{\log T}) \text{ when } \pi' = \text{KL-UCB}. \quad (20)$$

5.2 SAP with more than two actions

In the SAP with two sensors, only action 2 provides information about the losses. In the case with $K > 2$ sensors, by playing an action k , we can obtain information about the losses of all sensors $l < k$ by recursively applying the dominance condition between pair of sensors. Further, any information provided by action $k > 2$ is contained in that provided by all actions $k' \geq k$ —if action k is played in round t , then we observe predictions $\{\hat{Y}_t^1, \hat{Y}_t^2, \dots, \hat{Y}_t^i\}$ which includes the observed predictions of all actions $k' \leq i$. This side-observation can be represented by a directed graph $G^S = (V, E)$, where $|V| = K$ and $E = \{(i, j) : i1 < i \leq j \leq K\}$. Note that G^S has self loops for all nodes except for node 1. The nodes in G^S represents actions of the SAP




Figure 2: Side observation graph G^S

and an edge $(i, j) \in E$ implies that actions i provides information about action j . The side-observation graph for the SAP is shown in Figure (2).

Theorem 2 *Let the dominance condition (10) holds. Then SAP ψ with $K \geq 2$ is regret equivalent to a MAB with side-observation graph G^S .*

Remark 1 *Note that the some of mean values $\{\gamma_1 - \gamma_i - \sum_{j \leq i} c_j\}$ need not be positive. Since the stochastic bandit algorithms assume that reward lie in the interval $[0, 1]$, we can ensure positive means by setting distribution ν_k , to have mean $\gamma_1 - \gamma_i - \sum_{j < i} c_j + \sum_{k \leq K-1} c_k$. Note that mean of each arm is shifted by the same amount, which does not change the regret value. This recovers the SAP with $K = 2$ actions and Theorem 1 holds.*

Proposition 3 (SAP regret lower bound) *Let π be any policy on SAT with 2 sensors such that it pulls the suboptimal arm only sub polynomial many times, i.e., $\mathbb{E}[N_i^\psi(T)] = o(T^a)$ for all $a > 0$ and $i \neq i^*$. Then,*

$$\liminf_{T \rightarrow \infty} R_T^\psi(\pi) / \log T \geq \kappa \text{ where} \quad (21)$$

$$\begin{aligned} \kappa = \min_{\{w_i\}} \sum_{i \in [K]} (\mu_{i^*} - \mu_i) w_i \\ \text{subjected to } \sum_{j \leq i} w_j \geq 1/D(\mu_i + \sum_{j < i} c_j || \mu_{i^*}) \text{ for all } i \in [K] \\ w_i \geq 0 \text{ for all } i \in [K] \end{aligned} \quad (22)$$

Proposition 4 (K-SAT regret upper bound) *Let π' denote a policy on a K -armed stochastic bandit where mean of arm $i > 1$ is $\gamma_1 - \gamma_i - \sum_{j < i} c_j$ and arm 1 has a fixed reward of value zero, and the side-observation graph is G^S . Then, the regret of a policy $g_1(\pi)$ for the SAT problem obtained from mapping (35) is upper bounded as*

$$R_T^\psi(g(\pi)) \leq \mathcal{O}(\xi(G^S) \log T + K^2) \quad (23)$$

when $\pi' = \text{UCB} - \text{LP}$ *Bucapatnam et al. (2014).*

6 Extension to context based prediction

In this section we consider that the prediction errors depend on the context X_t , and in each round the learner can decide which action to apply based on X_t . Let $\gamma_i(X_t) = \Pr\{\hat{Y}_t^1 \neq \hat{Y}_t^2 | X_t\}$ for all $i \in [K]$. We refer to this setting as Contextual Sensor Acquisition Problem (CSAP) and denote it as $\psi_c = (K, \mathcal{A}, \mathcal{C}, (\gamma_i, c_i)_{i \in [K]})$.

Given $x \in \mathcal{C}$, let $L_t(a|x)$ denote the loss from action $a \in \mathcal{A}$ in round t . A policy on ϕ^c maps past history and current contextual information to an action. Let Π^{ψ_c} denote set of policies on ψ_c and for any policy $\pi \in \Pi^{\psi_c}$, let $\pi(x_t)$ denote the action selected when the context is x_t . For any sequence $\{x_t, y_t\}_{t \geq 0}$, the regret of a policy π is defined as:

$$R_T^{\phi^c}(\pi) = \sum_{t=1}^T \mathbb{E}[L_t(\pi(x_t)|x_t)] - \sum_{t=1}^T \min_{a \in \mathcal{A}} \mathbb{E}[L_t(a|x_t)]. \quad (24)$$

As earlier, the goal is to learn a policy that minimizes the expected regret, i.e., $\pi^* = \arg \min_{\pi \in \Pi^{\psi_c}} \mathbb{E}[R_T^{\psi_c}(\pi)]$.

In this section we focus on CSA-problem with two sensors and assume that sensor predictions errors are linear in the context. Specifically, we assume that there exists $\theta_1, \theta_2 \in \mathcal{R}^d$ such that $\gamma_1(x) = x' \theta_1$ and $\gamma_2(x) + c = x' \theta_2$ for all $x \in \mathcal{C}$, where x' denotes the transpose of x . By default all vectors are column vectors. In the following we establish that CSAP is regret equivalent to a stochastic linear bandits with varying decision sets. We first recall the stochastic linear bandit setup and relevant results.

Note: c is a fixed cost and does not depend on context. We are assuming that error rate of sensor 2 offset by c is a linear quantity. Another possibility is, we can assume that there exists a $x_0 \in \mathcal{C}$ such that $c = x_0' \theta_2$ and we have oracle access to x_0 . Then all the arguments hold.

7 Background on Stochastic Linear Bandits

In round t , the learner is given a decision set $D_t \subset \mathcal{R}^d$ from which he has to choose an action. For a choice $x_t \in D_t$, the learner receives a reward $r_t = x_t' \theta^* + \epsilon_t$, where $\theta^* \in \mathcal{R}^d$ is unknown and ϵ_t is random noise of zero mean. The learner's goal is to maximize the expected accumulated reward $\mathbb{E} \left[\sum_{t=1}^T r_t \right]$ over a period T . If the learner knows θ^* , his optimal strategy is to select $x_t^* = \arg \max_{x \in D_t} x' \theta^*$ in round t . The performance of any policy π that selects action x_t at time t is measured with respect to the optimal policy and is given by the expected regret as follows

$$R_T^L(\pi) = \sum (x_t^*)' \theta^* - \sum x_t' \theta^*. \quad (25)$$

The above setting, where actions sets can change in every round, is introduced in Abbasi-Yadkori et al. (2011) and is a more general setting than that studied in Dani et al. (2008); Rusmevichientong & Tsitsiklis (2010) where decision set is fixed. Further, the above setting also specializes the contextual bandit studied in Li et al. (2010). The authors in Abbasi-Yadkori et al. (2011) developed an 'optimism in the face of uncertainty linear bandit algorithm' (OFUL) that achieves $\mathcal{O}(d\sqrt{T})$ regret with high probability when the random noise is R -sub-Gaussian for some finite R . The performance of OFUL is significantly better than *ConfidenceBall*₂ Dani et al. (2008), *UncertaintyEllipsoid* Rusmevichientong & Tsitsiklis (2010) and *LinUCB* Li et al. (2010).

Theorem 3 Consider a CSA-problem with $K = 2$ sensors. Let \mathcal{C} be a bounded set and $\gamma_i(x) + c_i = x' \theta_i$ for $i = 1, 2$ for all $x \in \mathcal{C}$. Assume $x' \theta_1, x' \theta_2 \in [0, 1]$ for all $x \in \mathcal{C}$. Then, equivalent to a stochastic linear bandit.

References

- Abbasi-Yadkori, Yasin, Pál, Dávid, and Szepesvári, Csaba. Improved algorithms for linear stochastic bandits. In *Proceeding of Advances in Neural Information Processing Systems (NIPS)*, pp. 2312–2320, 2011.
- Audibert, J.-Y. and Bubeck, S. Regret bounds and minimax policies under partial monitoring. *Journal of Machine Learning Research*, 11:2635–2686, 2010.
- Auer, P., Nicholó-Cesa-Bianchi, and Fischer, Paul. Finite-time analysis of multiarmed bandit problem trade-offs. *Journal of Machine Learning*, 3:235–256, 2002.
- Buccapatnam, S., Eryilmaz, A., and Shroff, N. B. Stochastic bandits with side observation on networks. In *Proceeding of Sigmetrics*, 2014.
- Dani, V., Hayes, T. P., and Kakade, S. M. Stochastic linear optimization under bandit feedback. In *Proceeding of Conference on Learning Theory, COLT*, Helsinki, Finland, July 2008.
- Garivier, A. and Cappé, O. The kl-ucb algorithm for bounded stochastic bandits and beyond. In *Proceedings of the Annual Conference on Learning Theory (COLT)*, 2011.
- Hanna, J. Approximation to bayes risk in repeated plays. *Contributions to the Theory of Games*, 3: 97–139, 1957.
- Helmsboat, D. P., Littlestone, PN, and Long, P.M. Apple tasting. *Journal of Information and Computation*, 161(2):85–139, 2000.
- Lai, Tze Leung and Robbins, Herbert. Asymptotically efficient adaptive allocation rules. *Journal of Advances in applied mathematics*, 6(1):4–22, 1985.
- Li, L., Wei, C., Langford, J., and Schapire, R. E. A contextual-bandit approach to personalized news article recommendation. In *Proceeding of International Word Wide Web conference, WWW*, NC, USA, April 2010.
- Mackay, David. J. C. *Information Theory, Inference, and Learning Algorithms*. Cambridge University Press.
- Rusmevichientong, Paat and Tsitsiklis, John N. Linearly parameterized bandits. *Mathematics of Operations Research*, 35(2):395–411, 2010.
- Seldin, Y., Bartlett, P., Crammer, K., and Abbasi-Yadkori, Y. Prediction with limited advice and multiarmed bandits with paid observations. In *Proceeding of International Conference on Machine Learning, ICML*, pp. 208–287, 2014.
- Trapeznikov, K. and Saligrama, V. Supervised sequential classification under budget constraints. In *Proceeding of International Conference on Artificial Intelligence and Statistics, AISTATS*, pp. 235–242, 2013.
- Trapeznikov, K., Saligrama, V., and Castanon, D. A. Multistage classifier design. *Machine Learning Journal*, 39:1–24, 2014.
- Xu, Z., Kusner, M., Chen, M., and Weinberger, K. Q. Cost-sensitive tree of classifiers. In *Proceeding of International Conference on Machine Learning, ICML*, pp. 133–141, 2013.

A Proof of Theorem 1

Consider a 1-armed stochastic bandit problem where arm with constant reward has value c and the arm that gives stochastic reward has mean value $\gamma_1 - \gamma_2$. Given an arbitrary policy $\pi = (\pi_1, \pi_2, \dots, \pi_t)$ for the SAP, we obtain a policy for the bandit problem from π as follows: Let H_{t-1} denote the history, consisting of all arms played and the corresponding rewards, available to policy π_{t-1} till time $t - 2$. Let a_{t-1} denote the action selected by the bandit policy in round $t - 1$ and r_{t-1} the observed reward. Then, the next action a_t is obtained as follows:

$$a_t = \begin{cases} \pi_t(H_{t-1} \cup \{1, \emptyset\}) & \text{if } a_{t-1} = \text{fixed reward arm} \\ \pi_t(H_{t-1} \cup \{2, r_{t-1}\}) & \text{if } a_{t-1} = \text{stochastic arm} \end{cases} \quad (26)$$

Conversely, let $\pi' = \{\pi'_1, \pi'_2, \dots\}$ denote an arbitrary policy for the 1-armed bandit problem. we obtain a policy for the SAP as follows: Let H'_{t-1} denote the history, consisting of all actions played and feedback, available to policy π'_{t-1} till time $t - 1$. Let a'_{t-1} denote the action selected by the SAP policy in round $t - 1$ and observed feedback F_t . Then, the next action a'_t is obtained as follows:

$$a'_t = \begin{cases} \pi'_t(H'_{t-1} \cup \{1, c\}) & \text{if } a'_{t-1} = \text{action 1} \\ \pi'_t(H'_{t-1} \cup \{2, 1\{\hat{Y}_t^1 \neq \hat{Y}_t^2\}\}) & \text{if } a'_{t-1} = \text{actions 2.} \end{cases} \quad (27)$$

We next show that regret of π on the SAP is same as that of derived policy on the 1-armed bandit, and regret of π' on the 1-armed bandit is same as regret of the derived policy on SAP. We first argue that any policy on the SAP problem with 2 actions needs the information if whether the predictions of sensors match or not whenever action 2 is played. The following observation is straightforward.

Lemma 1 *Let dominance condition holds. Then, $\Pr\{\hat{Y}_t^1 \neq \hat{Y}_t^2\} = \gamma_1 - \gamma_2$.*

$$\Pr\{\hat{Y}_t^1 \neq \hat{Y}_t^2\} = \Pr\{\hat{Y}_t^1 = Y_t, \hat{Y}_t^2 \neq Y_t\} + \Pr\{\hat{Y}_t^2 = Y_t, \hat{Y}_t^1 \neq Y_t\} \quad (28)$$

$$= \Pr\{\hat{Y}_t^2 = Y_t, \hat{Y}_t^1 \neq Y_t\} \quad \text{from assumption (10)} \quad (29)$$

$$= \Pr\{\hat{Y}_t^1 \neq Y_t\} \Pr\{\hat{Y}_t^2 = Y_t | \hat{Y}_t^1 \neq Y_t\} \quad (30)$$

$$= \Pr\{\hat{Y}_t^1 \neq Y_t\} \left(1 - \Pr\{\hat{Y}_t^2 \neq Y_t | \hat{Y}_t^1 \neq Y_t\}\right) \quad (31)$$

$$= \Pr\{\hat{Y}_t^1 \neq Y_t\} \left(1 - \frac{\Pr\{\hat{Y}_t^2 \neq Y_t, \hat{Y}_t^1 \neq Y_t\}}{\Pr\{\hat{Y}_t^1 \neq Y_t\}}\right) \quad (32)$$

$$= \Pr\{\hat{Y}_t^1 \neq Y_t\} - \Pr\{\hat{Y}_t^2 \neq Y_t\} \quad \text{by contrapositive of (10)} \quad (33)$$

From Lemma 1, mean of the observations $Z_t := 1\{\hat{Y}_t^1 \neq \hat{Y}_t^2\}$ from action 2 in the SAP is a sufficient statistics to identify the optimal arm. Thus, any SAP only needs to know Z_t in each round, and Z_t are i.i.d with mean $\gamma_1 - \gamma_2$. Our mapping of policies is such that any policy for SAP (1-armed bandits) and the derived policy on the 1-armed bandit (SAP) play the sub-optimal arm same number of times. For the sake of simplicity assume that action 1 is optimal for SAP ($\gamma_1 > \gamma_2 + c$) and let a policy π on SAP plays it $N_1(T)$ number of times. Then, we have

$$R_T^\psi(\pi) = \Delta_i \mathbb{E}[N_1^\psi(T)] = (\gamma_1 - \gamma_2 - c) \mathbb{E}[N_1(T)]$$

Let $f(\pi)$ denote the policy for the 1-armed bandit obtained using the mapping (26). Now, for the 1-armed bandit, where the arm with stochastic rewards is optimal, we have

$$R_T^\phi(f(\pi)) = (\mu_2 - \mu_1) \mathbb{E}[N_1(T)] = (\gamma_1 - \gamma_2 - c) \mathbb{E}[N_1^\phi(T)]$$

Thus the regret of π on the SAP problem and that of $f(\pi)$ on the 1-armed bandit are the same. We can argue similarly for the other case.

B Proof of Theorem 2

Consider a K -armed stochastic bandit problem where rewards distribution ν_i has mean $\gamma_1 - \gamma_i - \sum_{j < i} c_j$ for all $i > 1$ and arm 1 gives a fixed reward of value 0. The arms have side-observation

structure defined by graph G^S . Given an arbitrary policy $\pi = (\pi_1, \pi_2, \dots, \pi_t)$ for the SAP, we obtain a policy for the bandit problem with side observation graph G^S from π as follows: Let H_{t-1} denote the history, consisting of all arms played and the corresponding rewards, available to policy π_{t-1} till time $t-2$. In round $t-1$, let a_{t-1} denote the arm selected by the bandit policy, r_{t-1} the corresponding reward and o_{t-1} the side-observation defined by graph G_S excluding that from the first arm. Then, the next action a_t is obtained as follows:

$$a_t = \begin{cases} \pi_t(H_{t-1} \cup \{1, \emptyset\}) & \text{if } a_{t-1} = \text{arm 1} \\ \pi_t(H_{t-1} \cup \{i, r_{t-1} \cup o_{t-1}\}) & \text{if } a_{t-1} = \text{arm } i \end{cases} \quad (34)$$

Conversely, let $\pi' = \{\pi'_1, \pi'_2, \dots\}$ denote an arbitrary policy for the K -armed bandit problem with side-observation graph. we obtain a policy the SAP as follows: Let H'_{t-1} denote the history, consisting of all actions played and feedback, available to policy π'_{t-1} till time $t-2$. Let a'_{t-1} denote the action selected by the SAP policy in round $t-1$ and observed feedback F_t . Then, the next action a'_t is obtained as follows:

$$a'_t = \begin{cases} \pi'_t(H'_{t-1} \cup \{1, 0\}) & \text{if } a'_{t-1} = \text{action 1} \\ \pi'_t(H'_{t-1} \cup \{i, \mathbf{1}\{\hat{Y}_t^1 \neq \hat{Y}_t^2\} \dots \mathbf{1}\{\hat{Y}_t^1 \neq \hat{Y}_t^i\}\}) & \text{if } a_{t-1} = \text{action } i. \end{cases} \quad (35)$$

We next show that regret of a policy π on the SAP problem is same as that of the policy derived from it for the K -armed bandit problem with side information graph G^S , and regret of π' on the K -armed bandit with side information graph G^S is same as that of the policy derived from it for the SAP.

Given a policy π for the SAP problem let $f_1(\pi)$ denote the policy obtained by the mapping defined in (34). The regret of policy π that plays actions i , $N_i(T)$ times is given by

$$R_T^\psi(\pi) = \sum_{i=1}^K \left[\left(\gamma_i + \sum_{j < i} c_j \right) - \left(\gamma_{i^*} + \sum_{j < i^*} c_j \right) \right] \mathbb{E}[N_i^\psi(T)] \quad (36)$$

$$(37)$$

Now, regret of regret policy $f_1(\pi)$ on the K -armed bandit problem with side information graph G^S

$$R_T^{\phi_G}(f_1(\pi)) = \sum_{i=1}^K \left[\left(\gamma_1 - \gamma_{i^*} - \sum_{j < i^*} c_j \right) - \left(\gamma_1 - \gamma_i - \sum_{j < i} c_j \right) \right] \mathbb{E}[N_i^{\phi_G}(T)] \quad (38)$$

which is same as $R_T^\phi(\pi)$. This concludes the proofs.

C Proof of Theorem 3

Let $\{x_t, y_t\}_{t \geq 0}$ be an arbitrary sequence of context-label pairs. Consider a stochastic linear bandit where $D_t = \{0, x_t\}$ is a decision set in round t . From the previous section, we know that given a context x , action 1 is optimal if $\gamma_1(x) - \gamma_2(x) - c < 0$, otherwise action 2 is optimal. Let $\theta := \theta_1 - \theta_2$, then it boils down to check if $x'\theta - c < 0$ for each context $x \in \mathcal{C}$.

For all t , define $\epsilon_t = \mathbf{1}\{\hat{Y}_t^1 \neq \hat{Y}_t^2\} - x'_t\theta$. Note that $\epsilon_t \in [0, 1]$ for all t , and since sensors do not have memory, they are conditionally independent given past contexts. Thus, $\{\epsilon_t\}_{t \geq 0}$ are conditionally R -sub-Gaussian for some finite R .

Given a policy π on a linear bandit we obtain next to play for the CSAP as follows: For each round t define $a_t \in \mathcal{C}$ and $r_t \in \{0, 1\}$ such that $a_t = 0$ and $r_t = 0$ if action 1 is played in that round, otherwise set $a_t = x_t$ and $r_t = \mathbf{1}\{\hat{y}_t^1 \neq \hat{y}_t^2\}$. Let $\mathcal{H}_t = \{(a_1, r_1) \dots (a_{t-1}, r_{t-1})\}$ denote the past actions and corresponding rewards observed till time $t-1$. In round t , after observing context x_t , we transfer $((a_{t-1}, r_{t-1}), D_t)$, where $D_t = \{0, x_t\}$. If π outputs $0 \in D_t$ as the optimal choice, we play action 1, otherwise we play action 2.

Conversely, suppose π' denote a policy for the CSAP problem we select action to play from decision set $D_t = \{0, x_t\}$ as follows. For each round t define $a'_t \in \{1, 2\}$ and $r'_t \in \mathcal{R}$ such that $a'_t = 1$ and $r'_t = \emptyset$ if 0 is played otherwise set $a'_t = 2$ and $r'_t = x'_t\theta^* + \epsilon_t$ if x_t is played. Let $\mathcal{H}'_t = \{(a'_1, r'_1) \dots (a'_{t-1}, r'_{t-1})\}$ denote the past actions and corresponding rewards observed till time $t-1$. In round t , after observing set D_t , we transfer $((a'_{t-1}, r'_{t-1}), x_t)$ to policy π' . If π outputs action 1 as the optimal choice, we play action 0, otherwise we play x_t .